

Analysis

February 24, 2023

1 Analysis

1.1 Contents

The notebook is laid out as follows:

- Loading the data
- Chart Data Exploration
- One Hit Wonders
- Last FM data
- Recommendations

I will also export it to PDF just in case it fails to render.

```
[1]: import pathlib
import sqlite3

import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns

pd.set_option("max_columns", 40)
```

1.2 Load Data

The data is stored in a sqlite database, comprised of four tables (linked either with `lastfm_id` or `artist_name`, `song_name`):

- charts: The UK chart data, scraped from the website.
- lastfm: The lastfm data for the tracks.
- lastfm_tags: The song tag data.
- lastfm_artists: The artists from lastfm.

This is primarily as sometimes I think it is better to work in SQL for certain types of analysis (in particular, it's hard to beat the expressiveness of window functions).

```
[2]: dataset_path = pathlib.Path("../data/dataset.sqlite")
conn = sqlite3.connect(dataset_path)
```

```
[3]: def sql(query: str) -> pd.DataFrame:
      """Helper function for querying the database."""
      return pd.read_sql_query(query, conn)
```

1.3 Sanity Checks

Below are some sanity checks for the data we're loading, to make sure it can be loaded and that it looks alright.

```
[4]: # Can we load the data?
charts = sql("SELECT * FROM charts ORDER BY week_ending DESC")
charts.head()

charts.to_excel("../data/charts_dump.xlsx", index=False)
```

```
[5]: # Check that the dates make sense.
query = """
SELECT MIN(year), MAX(year), COUNT(DISTINCT year), COUNT(DISTINCT chart_id),
       COUNT(*) FROM
(SELECT STRFTIME('%Y', week_ending) AS year, chart_id FROM charts) a
"""
sql(query)
```

```
[5]:  MIN(year) MAX(year)  COUNT(DISTINCT year)  COUNT(DISTINCT chart_id) \
0      1980      2010                31                1566

      COUNT(*)
0      148987
```

```
[6]: # 1566 * 100 != 148987 - which I think is because the charts didn't always have
      ↳ 100 entries.
query = """
SELECT chart_id, COUNT(*) AS cnt
FROM charts
GROUP BY chart_id
"""
sql(query)
```

```
[6]:      chart_id  cnt
0      7501-19800105  75
1      7501-19800112  75
2      7501-19800119  75
3      7501-19800126  75
4      7501-19800202  75
...      ...      ...
1561    7501-20091205  100
1562    7501-20091212  100
```

```

1563  7501-20091219  100
1564  7501-20091226  100
1565  7501-20100102  100

```

```
[1566 rows x 2 columns]
```

1.4 Chart Data Exploration

Below are charts and tables that I used to understand the dataset.

```

[7]: # Summary stats
sql(
    """
    SELECT
        COUNT(DISTINCT artist_name),
        COUNT(DISTINCT chart_id),
        COUNT(DISTINCT product_id),
        COUNT(DISTINCT song_name),
        COUNT(DISTINCT label_name)
    FROM charts
    """
).T

```

```

[7]:
COUNT(DISTINCT artist_name)  11832
COUNT(DISTINCT chart_id)    1566
COUNT(DISTINCT product_id)  32336
COUNT(DISTINCT song_name)   25642
COUNT(DISTINCT label_name)   2471

```

```

[8]: # Filtering out outliers with 100+ chart appearances
df = sql(
    """
    SELECT artist_name, COUNT(DISTINCT chart_id) AS chart_appearances,
    COUNT(DISTINCT product_id) AS unique_songs
    FROM charts
    GROUP BY artist_name
    HAVING chart_appearances < 100
    """
)


```

```

[9]: # What's the spread of artists by the number of appearances they've made in the
    charts?
df.chart_appearances.hist()


```

```
[9]: <AxesSubplot:>
```



```
[10]: # What's the spread of artists by the number of appearances they've made in the charts?  
df.unique_songs.hist()
```

```
[10]: <AxesSubplot:>
```



```
[11]: # How well do songs do?
df = sql(
    """
    SELECT artist_name, song_name, MIN(week_ending) AS first_week_ending,
    COUNT(DISTINCT chart_id) AS total_weeks
    FROM charts
    GROUP BY 1, 2
    """
)
df.sort_values(by="total_weeks", ascending=False).head(10)
```


```
[11]:
```

	artist_name	song_name	first_week_ending	\
25054	SNOW PATROL	CHASING CARS	2006-07-29 00:00:00	
19900	OASIS	WHATEVER	1994-12-31 00:00:00	
26634	TAKE THAT	RULE THE WORLD	2007-10-27 00:00:00	
19891	OASIS	SOME MIGHT SAY	1995-05-06 00:00:00	
14799	KILLERS	MR BRIGHTSIDE	2004-06-05 00:00:00	
19878	OASIS	CIGARETTES & ALCOHOL	1994-10-22 00:00:00	
10062	FLO RIDA FT T-PAIN	LOW	2008-02-16 00:00:00	
19423	NEW ORDER	BLUE MONDAY	1983-03-19 00:00:00	
19902	OASIS	WONDERWALL	1995-11-11 00:00:00	
1257	AMY WINEHOUSE	REHAB	2006-10-28 00:00:00	

	total_weeks
25054	137
19900	110
26634	92
19891	81
14799	81
19878	79
10062	75
19423	74
19902	74
1257	73


```
[12]: df.total_weeks.hist()
```

```
[12]: <AxesSubplot:>
```



```
[13]: # How many new entrants are there over time?
df = sql(
    """
    SELECT artist_name, MIN(week_ending) AS first_chart_week
    FROM charts
    GROUP BY 1
    """
)
new_artists = pd.to_datetime(df.first_chart_week).dt.year
# .sort_index().plot.bar(figsize=(12,6))
pd.cut(
    new_artists, bins=[1980, 1985, 1990, 1995, 2000, 2005, 2010], right=False
).value_counts().sort_index().plot.bar()
```

```
[13]: <AxesSubplot:>
```




```
[14]: # Can we understand how songs perform on a track over time?
# The drop off % might be useful for estimating longer term performance
weeks_on_chart = sql(
    """
    SELECT
        SUBSTR(STRFTIME('%Y', MIN(week_ending) OVER (PARTITION BY artist_name,
        ↪song_name)), 0, 4) AS decade,
        product_id,
        artist_name,
        song_name,
        1 + (julianday(week_ending) - julianday(MIN(week_ending) OVER (PARTITION BY
        ↪artist_name, song_name)))/7 AS weeks_on_chart,
        position
    FROM charts
    ORDER BY weeks_on_chart DESC
    """
)
time_on_chart = weeks_on_chart.groupby("product_id").weeks_on_chart.max()
time_on_chart = time_on_chart[time_on_chart == 12]
```

```

weeks_on_chart = weeks_on_chart[weeks_on_chart.product_id.isin(time_on_chart.
    ↪index)]
grouped = weeks_on_chart.groupby(["decade"]).apply(
    lambda g: g.groupby("weeks_on_chart").agg({"position": ["mean", "sem"]}
)
songs_by_decade = weeks_on_chart.groupby(["decade"]).product_id.nunique().
    ↪to_dict()
fig, ax = plt.subplots(figsize=(8, 6))
for decade, group in grouped.reset_index().groupby("decade"):
    group.plot.line(
        x="weeks_on_chart",
        y=("position", "mean"),
        figsize=(12, 6),
        xlim=[1, 12],
        legend=True,
        ax=ax,
        label=f"{decade}Os (n={songs_by_decade[decade]})",
    )
plt.gca().invert_yaxis()

```



1.5 One Hit Wonders

1.5.1 Identifying One Hit Wonders

I have defined one hit wonders as when an artist has the majority of their top 40 success concentrated around one song (I've picked 2/3 of hit weeks on the chart).

I have also filtered out some terms for things like "FT" in order to limit the number of collaborations in the charts.


```
[15]: query = """
SELECT * FROM (
    SELECT
        artist_name,
        song_name,
        first_week_ending,
        max_chart_run,
        peak_position,
        CAST(max_chart_run AS float) / SUM(max_chart_run) OVER (PARTITION BY
↪artist_name) share_of_weeks -- What % of weeks did this particular song
↪occupy?
    FROM (
        SELECT
            artist_name,
            song_name,
            MIN(week_ending) AS first_week_ending,
            MIN(peak_position) AS peak_position,
            MAX(weeks_on_chart) AS max_chart_run
        FROM
            charts
        GROUP BY 1, 2
        HAVING peak_position <= 40
    )
)
WHERE 1=1 -- Helper statement
    AND share_of_weeks >= 0.67 -- must be at least above 50% to achieve "one hit"
    -- Filter out collaborations
    AND artist_name NOT LIKE '%FT%'
    AND artist_name NOT LIKE '%FEATURING%'
    AND artist_name NOT LIKE '%&%'
    AND artist_name NOT LIKE '%AND%'
    AND artist_name NOT LIKE '%VS%'
    AND artist_name NOT LIKE '%/%'
ORDER BY 3 DESC, 4 DESC
"""
one_hit_artists = sql(query)
one_hit_artists.to_excel("../data/output_one_hit_artists.xlsx", index=False)
```

```
[16]: # Helper view for later joins + the lastfm script.
view = f"""
CREATE VIEW v_one_hit_wonders AS
{query}
"""
# view = conn.execute(view)
```

```
[17]: one_hit_wonders = sql("SELECT * FROM v_one_hit_wonders")
one_hit_artists["first_week_ending"] = pd.to_datetime(
```

```
one_hit_artists["first_week_ending"]
)
```

```
[18]: # Not often you get to look for bangers as part of your sense check
one_hit_artists.sort_values(by="max_chart_run", ascending=False).head(10)
```

```
[18]:
```


	artist_name	song_name	first_week_ending	max_chart_run	\
57	JASON MRAZ	I'M YOURS	2008-12-06	54	
158	FRAY	HOW TO SAVE A LIFE	2007-01-27	53	
281	BODYROCKERS	I LIKE THE WAY	2005-04-30	52	
170	GOO GOO DOLLS	IRIS/STAY WITH YOU	2006-10-21	51	
90	SAM SPARRO	BLACK & GOLD	2008-03-29	47	
266	DANIEL POWTER	BAD DAY	2005-08-06	45	
126	NEWTON FAULKNER	DREAM CATCH ME	2007-08-04	43	
975	TELETUBBIES	TELETUBBIES SAY EH-OH!	1997-12-13	41	
91	GABRIELLA CILMI	SWEET ABOUT ME	2008-03-15	41	
71	MADCON	BEGGIN'	2008-08-23	40	

	peak_position	share_of_weeks
57	11	1.000000
158	4	0.697368
281	3	1.000000
170	26	0.927273
90	2	1.000000
266	2	1.000000
126	7	1.000000
975	1	1.000000
91	6	0.759259
71	5	1.000000

1.5.2 Analysis


```
[19]: # Average runs on the chart rose dramatically from 2000-2010 - they might be
      ↪ better known.
one_hit_artists.groupby(
    one_hit_artists.first_week_ending.dt.year
).max_chart_run.mean().plot.line()
```

```
[19]: <AxesSubplot:xlabel='first_week_ending'>
```




```
[20]: # Can we measure "sucess" by longevity in the charts?
pareto_artists_by_num_weeks = (
    one_hit_artists.groupby("max_chart_run").size().div(len(one_hit_artists)).
    ↪cumsum()
)
ax = pareto_artists_by_num_weeks.plot.line(
    figsize=(6, 2),
    xlabel="Weeks on Chart",
    ylabel="Fraction of Hits",
    xlim=[1, 52],
    ylim=[0, 1],
)
percentile = pareto_artists_by_num_weeks[pareto_artists_by_num_weeks > 0.84].
    ↪index[0]
plt.axvline(percentile, color="k", linestyle="--")
plt.axhline(0.84, color="k", linestyle="--")
```

```
[20]: <matplotlib.lines.Line2D at 0x7f38c5e6e1c0>
```




```
[21]: # Success? Number of runs on the chart - majority aren't that successful (<5
      ↪ weeks on the chart, but some really linger on for up to a year on the chart)
      one_hit_artists.max_chart_run.hist(bins=20)
```

[21]: <AxesSubplot:>




```
[22]: # Do one hit artists typically tend to have other tracks that do well? - No
      one_hit_artists.share_of_weeks.hist(bins=20)
```

[22]: <AxesSubplot:>



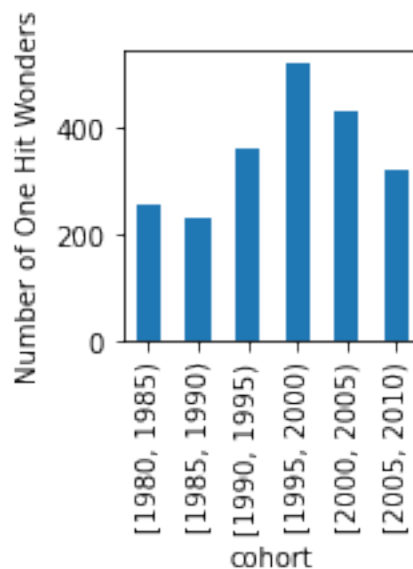
```
[23]: # Too detailed, but showcases how there's a lot of growth over the last 90s/00s.
one_hit_artists.first_week_ending.dt.year.value_counts().sort_index().plot.barh(
    figsize=(16, 8)
)
```

[23]: <AxesSubplot:>



```
[24]: # One hit wonders start to grow up until 2000 where it begins to peak.
# Q: who were people listening to post recession?
one_hit_wonders_by_cohort = (
    pd.cut(
        one_hit_artists.first_week_ending.dt.year,
        bins=[1980, 1985, 1990, 1995, 2000, 2005, 2010],
        right=False,
    )
    .value_counts()
    .to_frame(name="one_hit_wonders")
)
one_hit_wonders_by_cohort.sort_index().plot.bar(
    legend=False,
    xlabel="cohort",
    ylabel="Number of One Hit Wonders",
    figsize=(2, 2), # Presentation friendly
)
```

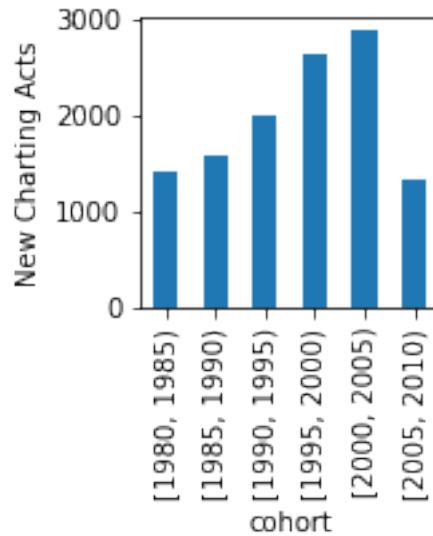
[24]: <AxesSubplot:xlabel='cohort', ylabel='Number of One Hit Wonders'>



```
[25]: # Can we see if there are lots of artists vying for the charts?
new_artists_by_cohort = (
    pd.cut(new_artists, bins=[1980, 1985, 1990, 1995, 2000, 2005, 2010],
    right=False)
    .value_counts()
    .to_frame(name="new_artists")
)
```

```
new_artists_by_cohort.sort_index().plot.bar(
    legend=False, xlabel="cohort", ylabel="New Charting Acts", figsize=(2, 2)
)
```

[25]: <AxesSubplot:xlabel='cohort', ylabel='New Charting Acts'>



```
[26]: # What tracks did well in the recession?
one_hit_artists[
    one_hit_artists.first_week_ending.dt.year.isin([2007, 2008, 2009])
].sort_values(by="max_chart_run", ascending=False).head(20)
```

```
[26]:
```

	artist_name	song_name	first_week_ending \
57	JASON MRAZ	I'M YOURS	2008-12-06
158	FRAY	HOW TO SAVE A LIFE	2007-01-27
90	SAM SPARRO	BLACK & GOLD	2008-03-29
126	NEWTON FAULKNER	DREAM CATCH ME	2007-08-04
91	GABRIELLA CILMI	SWEET ABOUT ME	2008-03-15
71	MADCON	BEGGIN'	2008-08-23
123	PLAIN WHITE T'S	HEY THERE DELILAH	2007-08-11
113	RAY PARKER JR	GHOSTBUSTERS	2007-11-10
67	MIA	PAPER PLANES	2008-09-13
133	FERGIE	BIG GIRLS DON'T CRY	2007-06-30
106	SOULJA BOY TELLEM	CRANK THAT (SOULJA BOY)	2007-12-08
124	ROBYN WITH KLEERUP	WITH EVERY HEARTBEAT	2007-08-11
64	GURU JOSH PROJECT	INFINITY 2008	2008-11-01
62	BELLAMY BROTHERS	LET YOUR LOVE FLOW	2008-11-08
49	METRO STATION	SHAKE IT	2009-03-07
37	DANIEL MERRIWEATHER	RED	2009-05-30

75	KID ROCK	ALL SUMMER LONG	2008-07-12
94	ONEREPUBLIC	STOP AND STARE	2008-02-23
42	VERONICAS	UNTOUCHED	2009-05-09
27	JLS	BEAT AGAIN	2009-07-25

	max_chart_run	peak_position	share_of_weeks
57	54	11	1.000000
158	53	4	0.697368
90	47	2	1.000000
126	43	7	1.000000
91	41	6	0.759259
71	40	5	1.000000
123	35	2	1.000000
113	35	2	0.777778
67	34	19	1.000000
133	33	2	0.702128
106	30	2	1.000000
124	28	1	1.000000
64	28	3	1.000000
62	26	7	1.000000
49	26	6	1.000000
37	26	5	0.787879
75	25	1	0.806452
94	25	4	1.000000
42	24	8	0.857143
27	24	1	0.750000

1.6 Last.FM Data

1.6.1 Data Exploration

```
[27]: # ~ 7 billion streams
sql("SELECT SUM(playcount) FROM lastfm")

[27]:      SUM(playcount)
0      723917022
```

```
[28]: # Slight tail off, possibly due to missed joins.
sql("SELECT COUNT(*) FROM lastfm")

[28]:      COUNT(*)
0      2016
```


```
[29]: # ~1 track has 3 labels
sql("SELECT COUNT(*) FROM lastfm_tags")
```



```
[29]: COUNT(*)  
0      7504
```


```
[30]: # What are some popular tags/genres?  
sql(  
    "SELECT tag, COUNT(*) FROM lastfm_tags GROUP BY 1 ORDER BY 2 DESC LIMIT 50"  
) .set_index("tag") .sort_values(by="COUNT(*)", ascending=True) .plot.  
    ↪ barh(figsize=(16, 8))
```

```
[30]: <AxesSubplot:ylabel='tag'>
```




```
[31]: # The average hit has between 10-100k streams  
sql("SELECT playcount FROM lastfm") .playcount .map(np.log10) .hist(bins=100)
```

```
[31]: <AxesSubplot:>
```




```
[32]: # Average hit has between 500-10,000 streams
sql("SELECT listeners FROM lastfm").listeners.map(np.log10).hist(bins=100)
```

```
[32]: <AxesSubplot:>
```




```
[33]: # Average artist has between 100k-1m streams
sql("SELECT playcount FROM lastfm_artists").playcount.map(np.log10).
      ↪ hist(bins=100)
```

[33]: <AxesSubplot:>



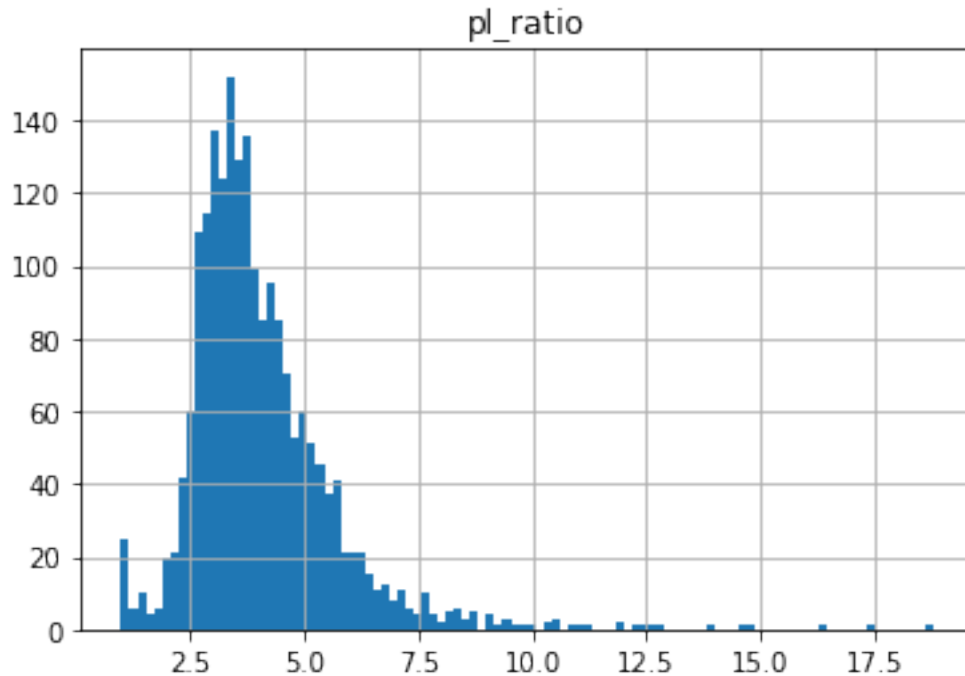
```
[34]: # Average artist has between 10k and 100k listeners
sql("SELECT listeners FROM lastfm_artists").listeners.map(np.log10).
      ↪ hist(bins=100)
```

[34]: <AxesSubplot:>



```
[35]: sql(
      "SELECT CAST(playcount AS FLOAT)/listeners AS pl_ratio FROM lastfm WHERE_
      ↪pl_ratio < 20"
    ).hist(bins=100)
```

```
[35]: array([[<AxesSubplot:title={'center':'pl_ratio'}>]], dtype=object)
```



```
[36]: pl_ratio_artists = sql(
        "SELECT CAST(playcount AS FLOAT)/listeners AS pl_ratio FROM lastfm_artists"
    )
    pl_ratio_artists.describe()
```

```
[36]:      pl_ratio
count  1814.000000
mean    9.152611
std     8.794331
min     2.346914
25%     4.513894
50%     6.227777
75%    10.192455
max    111.535817
```

1.6.2 Analysing songs based on streaming habits

```
[37]: # Join together the tracks and add some features such as listener stickiness_
      ↪(song_pl_ratio)
query = """
SELECT
    DISTINCT c.*,
    l.*,
    la.playcount AS artist_playcount,
```

```

    la.listeners AS artist_listeners,
    CAST(l.playcount AS FLOAT)/l.listeners AS song_pl_ratio,
    CAST(la.playcount AS FLOAT)/la.listeners AS artist_pl_ratio,
    CAST(l.listeners AS FLOAT)/la.listeners AS song_artist_listener_ratio,
    CAST(l.playcount AS FLOAT)/la.playcount AS song_artist_playcount_ratio
FROM lastfm l
JOIN link_lastfm_chart lc ON lc.lastfm_id = l.lastfm_id
JOIN lastfm_artists la ON la.artist_mbid = l.artist_mbid
JOIN v_one_hit_wonders c ON c.artist_name = lc.artist_name AND c.song_name = lc.
    ↳song_name
"""
ohw_stats = sql(query)
ohw_stats["first_week_ending"] = pd.to_datetime(ohw_stats.first_week_ending)

# sqlite annoyingly isn't as good an analytics database as I remember
# Add order of magnitude values for clearer reasoning
log10_cols = ["artist_playcount", "artist_listeners", "playcount", "listeners"]
for col in log10_cols:
    ohw_stats[f"log10_{col}"] = np.log10(ohw_stats[col])
ohw_stats["pct_split_playcount"] = (
    ohw_stats.sort_values(by="playcount", ascending=False)
    .playcount.cumsum()
    .div(ohw_stats.playcount.sum())
)
ohw_stats["playcount_rank"] = ohw_stats.playcount.rank(method="min",
    ↳ascending=False)

# Generate a year cohort to make grouping easier
ohw_stats["year_cohort"] = pd.cut(
    ohw_stats.first_week_ending.dt.year,
    bins=[0, 1980, 1985, 1990, 1995, 2000, 2005, 2010],
    right=False,
)[ohw_stats.index]

```

```

[38]: # Can determine weighting from this
ohw_stats[["year_cohort"]].value_counts(normalize=True).sort_index()

```

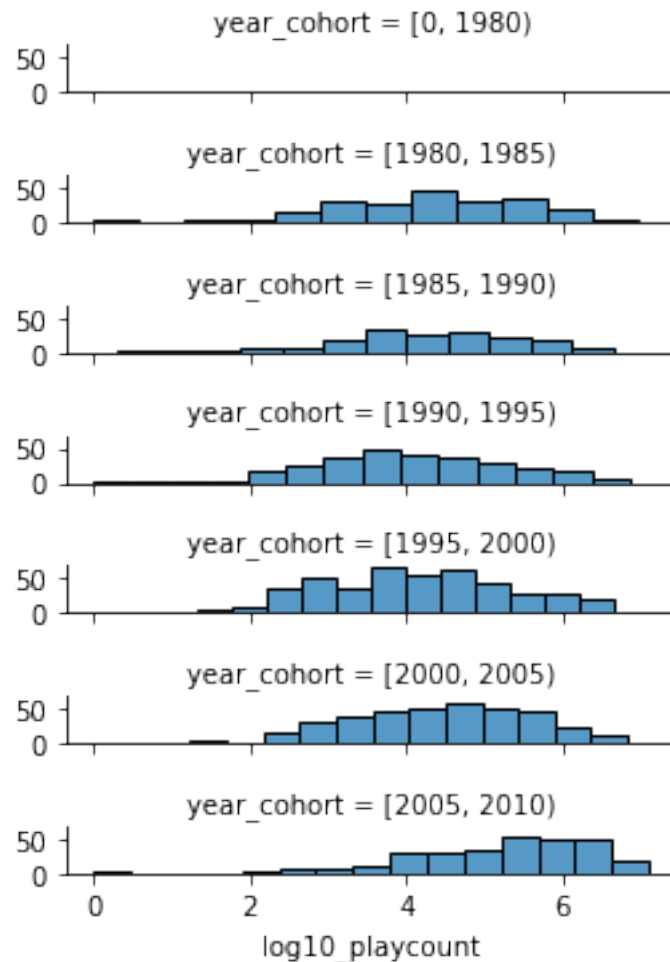
```

[38]: year_cohort
[0, 1980)      0.000000
[1980, 1985)   0.124728
[1985, 1990)   0.098039
[1990, 1995)   0.158497
[1995, 2000)   0.247821
[2000, 2005)   0.208061
[2005, 2010)   0.162854
dtype: float64

```


```
[39]: # Later charts (around advent of streaming) have more skew than earlier tracks
g = sns.FacetGrid(ohw_stats, row="year_cohort", aspect=5, height=0.75)
g.map(sns.histplot, "log10_playcount")
```

```
[39]: <seaborn.axisgrid.FacetGrid at 0x7f38e642cf40>
```



```
[40]: # Average is way higher
ohw_stats.groupby("year_cohort").playcount.mean().div(1_000_000).sort_index().
    plot.bar(
        ylabel="Average Streams (millions)", xlabel="Cohort"
    )
```

```
[40]: <AxesSubplot:xlabel='Cohort', ylabel='Average Streams (millions)'\>
```



1.6.3 Analysing Tag Data

```
[41]: tags = sql(f"""SELECT * FROM lastfm_tags""")
```

```
[42]: # Build a genres table with some well known genres - nb, some tracks overlap
genres = [
    "indie",
    "folk",
    "pop",
    "dance",
    "rock",
    "alternative",
    "metal",
    "classic rock",
]
genres_table = tags[tags.tag.isin(genres)][["lastfm_id", "tag"]]
genres_table = (
    pd.get_dummies(genres_table, prefix="genre", columns=["tag"])
    .groupby("lastfm_id")
    .sum()
```



```

        .reset_index()
    ) # Annoying bit of logic for uniqueness (one-to-many tags)
    genres_table

```

```

[42]:      lastfm_id  genre_alternative  genre_classic rock  genre_dance  \
0           4           0           0           1
1           8           0           0           0
2           9           0           0           0
3          10           0           0           1
4          11           0           0           0
...
1018       2008           1           0           0
1019       2009           0           0           1
1020       2010           0           0           0
1021       2011           0           0           0
1022       2013           0           0           1

```

```

      genre_folk  genre_indie  genre_metal  genre_pop  genre_rock
0           0           0           0           1           0
1           0           1           0           0           1
2           0           0           0           1           0
3           0           0           0           0           0
4           0           0           0           1           0
...
1018       0           0           0           0           1
1019       0           0           0           0           0
1020       0           1           0           1           0
1021       0           0           0           1           0
1022       0           0           0           0           0

```

[1023 rows x 9 columns]

```

[43]: # Build a time period table
time_periods = ["60s", "70s", "80s", "90s", "00s"]
time_table = tags[tags.tag.isin(time_periods)][["lastfm_id", "tag"]]
time_table = (
    pd.get_dummies(time_table, prefix="decade", columns=["tag"])
    .groupby("lastfm_id")
    .sum()
    .reset_index()
)
time_table

```

```

[43]:      lastfm_id  decade_00s  decade_60s  decade_70s  decade_80s  decade_90s
0           1           0           0           0           1           0
1           4           1           0           0           0           0
2           6           0           0           0           0           1

```


3	11	0	0	0	1	1
4	13	0	0	0	1	0
..
569	1990	0	0	0	0	1
570	2000	0	0	0	1	0
571	2001	0	0	0	1	0
572	2004	1	0	0	0	0
573	2013	0	0	0	0	1

[574 rows x 6 columns]

```
[44]: ohw_stats = ohw_stats.merge(time_table, on="lastfm_id", how="outer").merge(
      genres_table, on="lastfm_id", how="outer"
    )
ohw_stats = ohw_stats[ohw_stats.artist_name.notnull()]
```


```
[45]: ohw_stats.filter(like="genre").mean().sort_values().plot.barh()
```

[45]: <AxesSubplot:>



```
[46]: # Nostalgia factor? Why tag with "90s" etc otherwise?
ohw_stats.filter(like="decade").mean().plot.barh()
```


[46]: <AxesSubplot:>



```
[47]: # 80% of listens come from about 233 of the songs
# 50% of listens come from 79 songs
ohw_stats.plot.scatter(
    x="playcount_rank",
    y="pct_split_playcount",
    xlim=[-5, 1750],
    ylim=[0, 1.1],
    xlabel="Ranking of # Streams",
    ylabel="Fraction of all streams",
)
ohw_stats[ohw_stats.pct_split_playcount <= 0.8].sort_values(
    by="pct_split_playcount", ascending=False
).iloc[0].playcount_rank


# Super hits
# plt.axvline(79, color="r", linestyle="--")
plt.axvline(233, color="r", linestyle="--")
plt.axhline(0.8, color="r", linestyle="--")
# plt.axhline(0.5, color="r", linestyle="--")
```

```
[47]: <matplotlib.lines.Line2D at 0x7f38c44406a0>
```




```
[48]: # Max peak position doesn't really have much of a bearing
sns.scatterplot(x=ohw_stats.peak_position, y=np.log10(ohw_stats.playcount))
```

```
[48]: <AxesSubplot:xlabel='peak_position', ylabel='playcount'>
```




```
[49]: # Seems to have a switchpoint around the 20 week run mark (1/2 year), but it's
      ↪ so concentrated it's hard to tell if it's statistically significant
      sns.scatterplot(x=ohw_stats.max_chart_run, y=np.log10(ohw_stats.playcount))
```

```
[49]: <AxesSubplot:xlabel='max_chart_run', ylabel='playcount'>
```




```
[50]: # A measure of how concentrated an artist's plays are around one song
      ohw_stats.song_artist_playcount_ratio.hist(bins=100)
```

```
[50]: <AxesSubplot:>
```




```
[51]: # measure of how concentrated an artist's listeners are around their one hit  
ohw_stats.song_artist_listener_ratio.hist(bins=100)
```

```
[51]: <AxesSubplot:>
```




```
[52]: # Some artists have very sticky / concentrated fan bases
ohw_stats.artist_pl_ratio.plot.hist(bins=90, xlim=[2, 40])
```

```
[52]: <AxesSubplot:ylabel='Frequency'>
```




```
[53]: # Appears like most artists plays are not heavily concentrated in their one_
      ↪ hits - possibly a proxy for popularity?
ohw_stats.song_artist_playcount_ratio.hist(bins=100)
```

```
[53]: <AxesSubplot:>
```



```
[54]: # Looks like artists have more concentrated listeners around their major hits,
      ↪ maybe it's a measure of their catalogue's longer term popularity?
      ohw_stats.song_artist_listener_ratio.hist(bins=100)
```

```
[54]: <AxesSubplot:>
```




```
[55]: # Top tracks with high listeners are not really big artists
ohw_stats.sort_values(by="song_artist_playcount_ratio", ascending=False).
      head(10)
```

```
[55]:
```

	artist_name	song_name	first_week_ending	\
1228	PETER CETERA	GLORY OF LOVE	1986-08-02	
1274	PROPELLERHEADS	SPYBREAK!	1997-05-17	
575	FOUR BUCKETEERS	THE BUCKET OF WATER SONG	1980-05-03	
833	KATE WINSLET	WHAT IF	2001-12-08	
559	FERRY AID	LET IT BE	1987-04-04	
1721	TWISTED X	BORN IN ENGLAND	2004-06-19	
222	BOY MEETS GIRL	WAITING FOR A STAR TO FALL	1988-11-26	
1199	PARTNERS IN KRYME	TURTLE POWER	1990-07-21	
18	4-4-2	COME ON ENGLAND	2004-06-19	
265	CAST FROM CASUALTY	EVERLASTING LOVE	1998-03-14	

	max_chart_run	peak_position	share_of_weeks	lastfm_id	listeners	\
1228	13.0	3.0	1.0	902	172318.0	
1274	2.0	40.0	1.0	828	111518.0	
575	6.0	26.0	1.0	487	121.0	
833	17.0	6.0	1.0	708	57044.0	
559	7.0	1.0	1.0	242	4743.0	
1721	3.0	9.0	1.0	693	573.0	
222	14.0	9.0	1.0	877	110650.0	
1199	10.0	1.0	1.0	1498	7622.0	
18	5.0	2.0	1.0	5	1223.0	
265	9.0	5.0	1.0	65	55.0	

	playcount	duration	mbid	\
1228	699315.0	260000.0	09970bad-8cfe-4c74-a940-be7fd13fa5f7	
1274	377752.0	420000.0	ed0c9373-6ffa-4c60-83d4-9b3b673c10c7	
575	307.0	172000.0	None	
833	287372.0	244000.0	af45aab5-1c23-4ef6-acb1-e1ebe73c3e03	
559	13265.0	NaN	1138308e-d67d-4b19-b4e0-c85382c2ec29	
1721	1857.0	180000.0	52acd439-d243-4ff7-9f25-e240a51d6bbf	
222	542361.0	321000.0	80214d59-9f71-4cee-b3fb-f677720eaa7b	
1199	28986.0	230000.0	00db40ea-7874-4a4f-8c2f-96cbc1988115	
18	4241.0	223000.0	553af8f8-75f6-444f-a658-997682893a4d	
265	187.0	171000.0	None	

	artist_mbid	artist_playcount	\
1228	cdcd53c9-f37c-48cc-b7e3-dfe34be22e92	31095.0	
1274	cbaafb20-eb74-421a-85f1-2bb6341aad23	40143.0	
575	5270d69b-bf85-4399-949e-3a971bf2e981	311.0	
833	ceb05831-03e8-4605-904d-894ee0492d00	307011.0	
559	590233d0-5ab2-4a67-8f43-a647c2100bdc	14261.0	

1721	a330ac7f-4722-465f-938f-7c7a703219f0	2027.0
222	1752572e-2179-4507-9214-b29a6f2d7888	592946.0
1199	7e3b1e89-9dc4-44ff-bb47-21be91a3dd72	32067.0
18	e4a47762-3b26-4263-a756-6fd2c3e425a8	4740.0
265	33542e9c-48af-4b98-8fbd-bda7383881c3	210.0

	artist_listeners	song_pl_ratio	artist_pl_ratio	\
1228	7084.0	4.058282	4.389469	
1274	11735.0	3.387363	3.420793	
575	122.0	2.537190	2.549180	
833	60293.0	5.037725	5.091984	
559	4924.0	2.796753	2.896223	
1721	625.0	3.240838	3.243200	
222	114253.0	4.901591	5.189763	
1199	8221.0	3.802939	3.900620	
18	1341.0	3.467702	3.534676	
265	59.0	3.400000	3.559322	

	song_artist_listener_ratio	song_artist_playcount_ratio	\
1228	24.324958	22.489629	
1274	9.503025	9.410159	
575	0.991803	0.987138	
833	0.946113	0.936032	
559	0.963241	0.930159	
1721	0.916800	0.916132	
222	0.968465	0.914689	
1199	0.927138	0.903920	
18	0.912006	0.894726	
265	0.932203	0.890476	

	log10_artist_playcount	log10_artist_listeners	log10_playcount	\
1228	4.492691	3.850279	5.844673	
1274	4.603610	4.069483	5.577207	
575	2.492760	2.086360	2.487138	
833	5.487154	4.780267	5.458444	
559	4.154150	3.692318	4.122707	
1721	3.306854	2.795880	3.268812	
222	5.773015	5.057868	5.734288	
1199	4.506058	3.914925	4.462188	
18	3.675778	3.127429	3.627468	
265	2.322219	1.770852	2.271842	

	log10_listeners	pct_split_playcount	playcount_rank	year_cohort	\
1228	5.236331	0.817518	251.0	[1985, 1990)	
1274	5.047345	0.889168	350.0	[1995, 2000)	
575	2.082785	0.999977	1714.0	[1980, 1985)	
833	4.756210	0.916507	410.0	[2000, 2005)	

559	3.676053	0.996425	1088.0	[1985, 1990)
1721	2.758155	0.999684	1480.0	[2000, 2005)
222	5.043951	0.851879	291.0	[1985, 1990)
1199	3.882069	0.991306	907.0	[1990, 1995)
18	3.087426	0.999085	1333.0	[2000, 2005)
265	1.740363	0.999991	1756.0	[1995, 2000)

	decade_00s	decade_60s	decade_70s	decade_80s	decade_90s	\
1228	0.0	0.0	0.0	1.0	0.0	
1274	NaN	NaN	NaN	NaN	NaN	
575	NaN	NaN	NaN	NaN	NaN	
833	NaN	NaN	NaN	NaN	NaN	
559	NaN	NaN	NaN	NaN	NaN	
1721	NaN	NaN	NaN	NaN	NaN	
222	0.0	0.0	0.0	1.0	0.0	
1199	0.0	0.0	0.0	1.0	0.0	
18	NaN	NaN	NaN	NaN	NaN	
265	NaN	NaN	NaN	NaN	NaN	

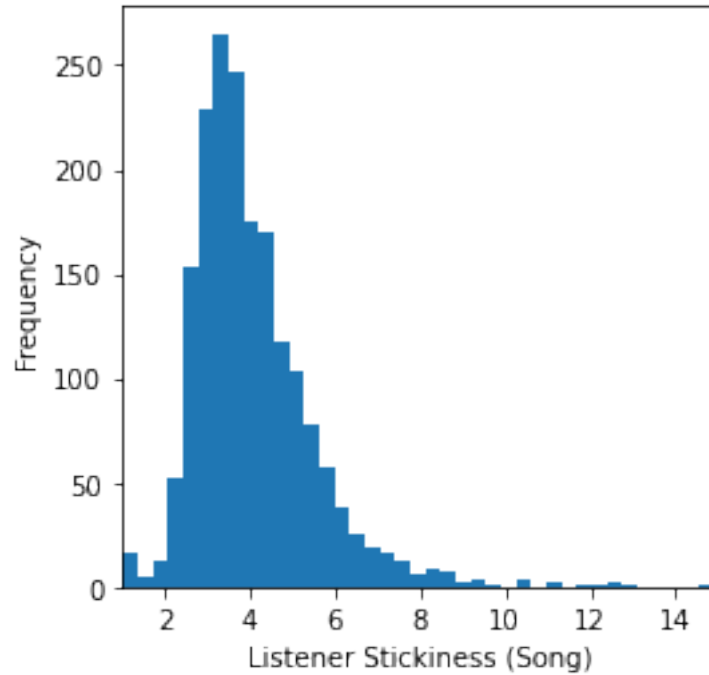
	genre_alternative	genre_classic rock	genre_dance	genre_folk	\
1228	0.0	0.0	0.0	0.0	
1274	0.0	0.0	1.0	0.0	
575	NaN	NaN	NaN	NaN	
833	0.0	0.0	0.0	0.0	
559	NaN	NaN	NaN	NaN	
1721	NaN	NaN	NaN	NaN	
222	0.0	0.0	0.0	0.0	
1199	NaN	NaN	NaN	NaN	
18	NaN	NaN	NaN	NaN	
265	NaN	NaN	NaN	NaN	

	genre_indie	genre_metal	genre_pop	genre_rock
1228	0.0	0.0	1.0	0.0
1274	0.0	0.0	0.0	0.0
575	NaN	NaN	NaN	NaN
833	0.0	0.0	1.0	0.0
559	NaN	NaN	NaN	NaN
1721	NaN	NaN	NaN	NaN
222	0.0	0.0	1.0	0.0
1199	NaN	NaN	NaN	NaN
18	NaN	NaN	NaN	NaN
265	NaN	NaN	NaN	NaN

```
[56]: # For artists, this is effectively a measure of how concentrated their fan base
      ↪ is in tandem with the play concentration
a = ohw_stats.song_pl_ratio.plot.hist(
    bins=90, xlim=[1, 15], xlabel="Listener Stickiness (Song)", figsize=(4, 4)
```


```
)  
a.set_xlabel("Listener Stickiness (Song)")
```

```
[56]: Text(0.5, 0, 'Listener Stickiness (Song)')
```




```
[57]: a = ohw_stats[  
    ohw_stats.song_artist_playcount_ratio <= 1  
].song_artist_playcount_ratio.plot.hist(  
    xlabel="Fraction of song plays vs total", figsize=(4, 4)  
)  
a.set_xlabel("Fraction of song plays vs total")
```

```
[57]: Text(0.5, 0, 'Fraction of song plays vs total')
```



```
[58]: # For artists, this is effectively a measure of how concentrated their fan base is
ohw_stats.artist_pl_ratio.hist(bins=25)
print(ohw_stats.artist_pl_ratio.mean())
```

9.184217837778288



```
[59]: ## Some feature engineering to get a feeling for how the tracks fare by
## May not include in final report, one big question that it raises is "when
↳were the majority of streams?" otherwise could hugely bias the dataset
## towards the more recent tracks.

import datetime

def years_since_spotify(row: pd.Timestamp) -> float:
    """Computes how many listens per year a track has had since the advent of
    ↳spotify."""
    end_date = pd.Timestamp(2021, 8, 6)
    reference_date = pd.Timestamp(2006, 4, 23)

    # If the song came out after spotify, then count that directly.
    if row > reference_date:
        reference_date = row
    return (end_date - reference_date).days / 365.25

ohw_stats["years_since_spotify"] = ohw_stats.first_week_ending.
    ↳map(years_since_spotify)
ohw_stats["annual_song_streams_since_spotify"] = (
    ohw_stats.playcount / ohw_stats.years_since_spotify
)
ohw_stats["annual_artist_streams_since_spotify"] = (
    ohw_stats.artist_playcount / ohw_stats.years_since_spotify
)

# https://www.indiemusicacademy.com/blog/music-royalties-explained
ohw_stats["gbp_annual_song_royalty_est"] = (
    ohw_stats.annual_song_streams_since_spotify * 0.066
)


[60]: # Come back if time - potentially too noisy with poor streaming data
# https://www.statista.com/statistics/885750/
↳average-risk-free-rate-united-kingdom/
risk_free_rate = 0.013
def npv(years: int, risk_free_rate: float, coupon: int) -> float:
    result = []
    for year in range(years):
        result.append(coupon / (1 + risk_free_rate) ** year)
    return np.sum(result)

# Below - bit dubious w/o trend data
```

```
# ohw_stats["gbp_annual_song_royalty_npv_5yr"] = ohw_stats.
↳gbp_annual_song_royalty_est.map(lambda r: npv(5, risk_free_rate, r))
# ohw_stats["gbp_annual_song_royalty_npv_10yr"] = ohw_stats.
↳gbp_annual_song_royalty_est.map(lambda r: npv(10, risk_free_rate, r))
# ohw_stats["gbp_annual_song_royalty_5yr_price_at_irr"] = ohw_stats.
↳gbp_annual_song_royalty_est.map(lambda r: npv(5, 0.1, r))
```


```
[61]: ohw_stats.annual_song_streams_since_spotify.map(np.log10).plot.hist()
```

```
[61]: <AxesSubplot:ylabel='Frequency'>
```



```
[62]: ohw_stats.annual_artist_streams_since_spotify.map(np.log10).plot.hist()
```

```
[62]: <AxesSubplot:ylabel='Frequency'>
```



1.6.4 Clustering Analysis

I use AffinityPropagation with a small feature set in order to group together songs with similar properties. This is a proof of concept to highlight why data science can be useful for performing fast target scanning / comp. analysis.

```
[63]: # Some extra stuff - quick and dirty clustering to see if we can segment the
      ↪ ohw market
      from sklearn.cluster import AffinityPropagation

      model = AffinityPropagation(damping=0.9)
      data = ohw_stats[
          [
              "song_artist_listener_ratio",
              "song_pl_ratio",
              "artist_pl_ratio",
              "peak_position",
              "share_of_weeks",
          ]
      ]
      model.fit(data)
      ohw_stats["cluster"] = model.predict(data)
```


```
/home/niall/Documents/Projects/[redacted]/assignment/venv/lib/python3.8/site-
packages/sklearn/cluster/_affinity_propagation.py:148: FutureWarning:
'random_state' has been introduced in 0.23. It will be set to None starting from
```


1.0 (renaming of 0.25) which means that results will differ at every function call. Set 'random_state' to None to silence this warning, or to 0 to keep the behavior of versions <0.23.

```
warnings.warn(
```

```
[64]: # Large skew - cluster 17 seems like a "catch all"
ohw_stats.cluster.value_counts().plot.barh(figsize=(16, 8))
```


```
[64]: <AxesSubplot:>
```



1.6.5 Searching for correlations


```
[65]: # Older songs don't appear to offer much of a leg up on newer songs for total
      ↪ streams
sns.scatterplot(x=ohw_stats.first_week_ending.dt.year, y=ohw_stats.
      ↪ log10_playcount)
```

```
[65]: <AxesSubplot:xlabel='first_week_ending', ylabel='log10_playcount'>
```




```
[66]: # Older songs don't appear to offer much of a leg up on newer songs for total
      ↪ streams
sns.scatterplot(
    x=ohw_stats.first_week_ending.dt.year, y=ohw_stats.log10_artist_playcount
)
```

```
[66]: <AxesSubplot:xlabel='first_week_ending', ylabel='log10_artist_playcount'>
```




```
[67]: # Nor for listeners except on average
sns.scatterplot(x=ohw_stats.first_week_ending.dt.year, y=ohw_stats.
↳ log10_listeners)
```

```
[67]: <AxesSubplot:xlabel='first_week_ending', ylabel='log10_listeners'>
```




```
[68]: # Nor for listeners, mean number of listeners is higher
sns.scatterplot(
    x=ohw_stats.first_week_ending.dt.year, y=ohw_stats.log10_artist_listeners
)
```

```
[68]: <AxesSubplot:xlabel='first_week_ending', ylabel='log10_artist_listeners'>
```




```
[69]: # The most listens the hit has, the better a predictor of the artist's overall_
      ↳ listens?
      sns.scatterplot(x=ohw_stats.log10_artist_playcount, y=ohw_stats.log10_playcount)

[69]: <AxesSubplot:xlabel='log10_artist_playcount', ylabel='log10_playcount'>
```




```
[70]: # Popular songs help artists with their overall portfolio
sns.scatterplot(x=ohw_stats.log10_artist_listeners, y=ohw_stats.log10_listeners)
```

```
[70]: <AxesSubplot:xlabel='log10_artist_listeners', ylabel='log10_listeners'>
```



```
[71]: sns.scatterplot(x=ohw_stats.log10_artist_listeners, y=ohw_stats.song_pl_ratio)
```

```
[71]: <AxesSubplot:xlabel='log10_artist_listeners', ylabel='song_pl_ratio'>
```



1.7 Recommendations

In order to build a stable portfolio that will generate revenue, I focus on the following criteria:

- Top half of listens (do people want to listen to the songs?)
- Avoiding potentially very pricy top tier songs
- Songs which performed well in the charts (3+ months)
- Artists who aren't too niche (as measured by playcount)

Plus some smaller data filtering steps (to get rid of some datapoints that I didn't think made sense, e.g. Jimi Hendrix as a one hit wonder).

```
[72]: mask = ohw_stats.pct_split_playcount <= 0.5  # pick popular songs
mask &= (
    ohw_stats.pct_split_playcount > 0.1
) # but not so popular that they'll be overpriced
mask &= (
    ohw_stats.duration < 4.5 * 60 * 1000
) # pick shorter songs (just to filter out the nicher classic rock stuff)
```

```

mask &= (
    ohw_stats.max_chart_run >= 12
) # pick songs that had more than 3 months of airtime on the charts
mask &= ohw_stats.song_artist_playcount_ratio <= 0.5
mask &= ohw_stats.decade_60s != 1
candidates = ohw_stats[mask]
candidates.sort_values(
    by=["year_cohort", "song_pl_ratio"], ascending=False
).sort_values(by="song_pl_ratio", ascending=False)

candidates.to_excel("../data/output_shortlist_review.xlsx", index=False)

```

[73]: *## I reviewed by eye to identify the songs fit for portfolio.*

```

selection_ids = [
    811, # Best PL for 80s,
    1456, # Best PL for 90s,
    252, # Best PL for early 00s,
    1755, # Third best PL for early 00s - different genre to best so I think
    ↪ that balances out better
    1605, # Best PL for late 00s
]

def format_output_table(row: pd.Series) -> pd.Series:
    return pd.Series(
        [
            row.lastfm_id,
            row.artist_name,
            row.song_name,
            tags[tags.lastfm_id == row.lastfm_id].tag.tolist(),
            row.first_week_ending.year,
            row.playcount,
            row.song_pl_ratio,
            row.song_artist_playcount_ratio,
            row.cluster,
        ],
        index=[
            "lastfm_id",
            "artist",
            "song",
            "tags",
            "year",
            "total_plays",
            "listener_stickiness",
            "listening_concentration",
            "cluster",
        ]
    )

```



```

    ],
)

final_candidates = (
    candidates.loc[selection_lds]
    .apply(format_output_table, axis=1)
    .sort_values(by="listener_stickiness", ascending=False)
)
final_candidates

```

```

[73]:      lastfm_id      artist      song \
1605      1262      TEMPER TRAP      SWEET DISPOSITION
811      1429      JOURNEY      DON'T STOP BELIEVIN'
252      723      CAESARS      JERK IT OUT
1456      355      SIXPENCE NONE THE RICHER      KISS ME
1755      1997      VANESSA CARLTON      A THOUSAND MILES

                                     tags year total_plays \
1605  [indie, indie rock, 500 Days of Summer, austra...  2009    9455762.0
811      [classic rock, 80s, rock, journey, scrubs]  1982    8982453.0
252      [indie, rock, alternative, indie rock, swedish]  2003    3002434.0
1456      [pop, 90s, female vocalists, Love, rock]  1999    4677837.0
1755  [pop, female vocalists, piano, vanessa carlton...  2002    3722367.0

listener_stickiness  listening_concentration  cluster
1605              9.492184              0.333081      30
811              6.970458              0.279932      23
252              6.476597              0.390669      36
1456              5.934677              0.460359      18
1755              5.552275              0.306132      36

```

1.7.1 Using the cluster model to construct a similar portfolio automatically

```

[74]: # Demo of fast lookup based on clustering
mask = ohw_stats.cluster.isin(
    final_candidates.cluster
) # Same clusters - some songs may be in a "unicluster", which is obviously a
↳drawback for the proof-of-concept
mask &= ~ohw_stats.lastfm_id.isin(
    final_candidates.lastfm_id
) # Don't pick the songs we know
mask &= ohw_stats.pct_split_playcount <= 0.5 # still focus on popular songs
mask &= ohw_stats.decade_60s != 1
ohw_stats[mask].sort_values(by="song_pl_ratio", ascending=False).groupby(
    "year_cohort"
).first().apply(format_output_table, axis=1).dropna().sort_values(

```

```
by="listener_stickiness"
)
```

```
[74]:          lastfm_id          artist \
year_cohort
[1995, 2000)      1457.0          NEW RADICALS
[2000, 2005)       404.0          JUNIOR SENIOR
[1980, 1985)     1676.0  DEXY'S MIDNIGHT RUNNERS WITH THE EMERALD EXPRESS
[1985, 1990)     1201.0          TRACY CHAPMAN
[2005, 2010)       53.0          JOSE GONZALEZ
```

```
          song \
year_cohort
[1995, 2000)  YOU GET WHAT YOU GIVE
[2000, 2005)      MOVE YOUR FEET
[1980, 1985)      COME ON EILEEN
[1985, 1990)      FAST CAR
[2005, 2010)      HEARTBEATS
```

```
          tags      year \
year_cohort
[1995, 2000)  [90s, rock, alternative, pop, alternative rock]  1999.0
[2000, 2005)      [dance, pop, electronic, fun, happy]  2003.0
[1980, 1985)      [80s, pop, new wave, british, rock]  1982.0
[1985, 1990)  [folk, acoustic, female vocalists, 80s, singer...  1988.0
[2005, 2010)  [acoustic, chillout, indie, Mellow, singer-son...  2006.0
```

```
          total_plays  listener_stickiness  listening_concentration \
year_cohort
[1995, 2000)      3231926.0          5.478704          0.600575
[2000, 2005)      2342399.0          5.510801          0.462756
[1980, 1985)      4271605.0          5.734545          0.662270
[1985, 1990)      4496842.0          6.739002          0.212117
[2005, 2010)      9539488.0          9.309897          0.208667
```

```
          cluster
year_cohort
[1995, 2000)      18.0
[2000, 2005)      18.0
[1980, 1985)      18.0
[1985, 1990)      23.0
[2005, 2010)      30.0
```


```
[75]: # Save the data for later
ohw_stats.to_excel("../data/output_one_hit_wonders_with_statistics.xlsx",
    ↪index=False)
sql("SELECT * FROM lastfm").to_excel("../data/output_lastfm.xlsx", index=False)
```

1.7.2 Some more tag analysis

```
[76]: # Are some genres more sticky than others?
pop_rock = ohw_stats.merge(
    sql("SELECT * FROM lastfm_tags WHERE tag in ('rock', 'dance', 'pop')"),
    on="lastfm_id",
)
```


```
[77]: pop_rock.groupby(
    ["tag", pop_rock.first_week_ending.dt.year]
).size().reset_index().pivot(index="first_week_ending", columns="tag",
    values=0).fillna(
    0
).plot.barh(
    figsize=(16, 8)
)
```

```
[77]: <AxesSubplot:ylabel='first_week_ending'>
```



```
[78]: fig, ax = plt.subplots()
pop_rock.groupby("tag").song_pl_ratio.plot.kde(
    legend=True, xlim=[2, 10], figsize=(4, 4)
)
ax.set_xlabel("Listener Stickiness")
```

```
[78]: Text(0.5, 0, 'Listener Stickiness')
```



```
[79]: pop_rock.groupby("tag").artist_pl_ratio.plot.kde(legend=True)
```

```
[79]: tag
dance    AxesSubplot(0.125,0.125;0.775x0.755)
pop      AxesSubplot(0.125,0.125;0.775x0.755)
rock     AxesSubplot(0.125,0.125;0.775x0.755)
Name: artist_pl_ratio, dtype: object
```

